

# Introduction to CUDA and OpenCL final project: Linear algebra toolkit

Mateusz Kaleta

February 2020

## 1 Introduction

The aim of this project was to create a toolkit application for linear algebra operations. Four operations were considered, namely: matrix addition/subtraction, multiplication and finding inverse using Gauss-Jordan elimination method. Each operation was implemented both for GPU and CPU in order to check the eventual speedup of GPU code. The application has an object oriented structure with different classes managing input/output and delegating tasks specified by the user.

## 2 Application structure

The application was implemented in C++ with the CUDA specific elements and libraries. It is build upon two classes.

### 2.1 Matrix class

This class provides the simple representation of a matrix and serves as an interface between user and the program. It stores matrix elements in one dimensional array, which is convenient for computation. It also provides methods to read matrices from text file, print out matrix elements to standard output etc.

### 2.2 Toolkit classes

Two toolkit classes: CPUToolkit and GPUSToolkit, which derive from abstract class Toolkit were implemented. Each class provides its own implementation (CPU and or GPU-style, respectively) of the four considered matrix operations.

## 3 Launching application

After compilation (e.g. make), the application can be launched command line with the following command, specifying different configurations:

```
./Toolkit [gpu—cpu] [add—subtract—multiply—inverse] [normal—performance]
```

where:

- first parameter specifies if we want to use GPU or CPU code,
- second parameter specifies what operation to perform,
- last parameter tell if we want to use "normal" mode: i.e. read matrices from text file and compute the result, or to launch an application in 'performance evaluation mode', i.e. perform the same operation for varying matrix sizes, to see how it influences the computation time.

## 4 Operations

### 4.1 Matrix addition/subtraction

The most basic operation implemented is matrix addition/subtraction:

$$\mathbf{C} = \mathbf{A} + \mathbf{B}, \quad (1)$$

which can easily be implemented in usual CPU style. However this requires iterating over every element of the matrix, and performing one operation at a time. The equivalent GPU code delegates the tasks to many cores, so that one core is responsible for single matrix element.

### 4.2 Matrix multiplication

Considering matrix multiplication, every element  $C_{ij}$  in the resulting matrix  $\mathbf{C}$  requires the computation of dot product of the corresponding row vector  $A_i$  and column vector  $B_j$ . Since every of this operation is independent, the parallel code can also be implemented. In this application the most standard CUDA-style matrix multiplication was implemented, where one core calculates one element of resulting matrix.




Figure 1: Visualisation of matrix multiplication with GPU.

Source: <https://www.quantstart.com/articles/Matrix-Matrix-Multiplication-on-the-GPU-with-Nvidia-CUDA/>

### 4.3 Matrix inversion: Gauss-Jordan elimination

To calculate the matrix inverse, the Gauss-Jordan elimination algorithm was implemented. Suppose we want to find an inverse of matrix  $\mathbf{A}_{n \times n}$ . We first create a corresponding identity matrix  $\mathbf{I}_{n \times n}$  of the same size. Then, by performing elementary operations on both  $\mathbf{A}_{n \times n}$  and  $\mathbf{I}_{n \times n}$  we aim to transform  $\mathbf{A}_{n \times n}$  into identity, which also turns  $\mathbf{I}$  into  $\mathbf{A}^{-1}$ :

$$\mathbf{A} \rightarrow \mathbf{I} = \mathbf{A}^{-1} \mathbf{A} \quad (2)$$

$$\mathbf{I} \rightarrow \mathbf{A}^{-1} \mathbf{I} = \mathbf{A}^{-1} \quad (3)$$

This operation consist of few steps. The corresponding pseudo-code is presented below.

For every row of  $\mathbf{A}$  and  $\mathbf{I}$ :

- If diagonal element  $A_{ii} = 0$ , find row element row k, such that  $A_{ki} \neq 0$  and add it to row i. Do the same for matrix **I**
- Divide elements in row i by diagonal element  $A_{ii}$ :  
 $A_{ij} = A_{ij}/A_{ii}$ , for  $i \neq j$   
 $I_{ij} = I_{ij}/A_{ii}$ .
- performed row reduction, i.e, subtract the current row multiplied by a scalar from all other rows in **A** and **I**.

## 5 Results

### 5.1 Computation time vs. matrix size

The GPU speedup was tested on three operations: addition (since subtraction is computationally of the same complexity), multiplication and Gauss-Jordan elimination. The computation time vs. matrix size is presented on the figures below.

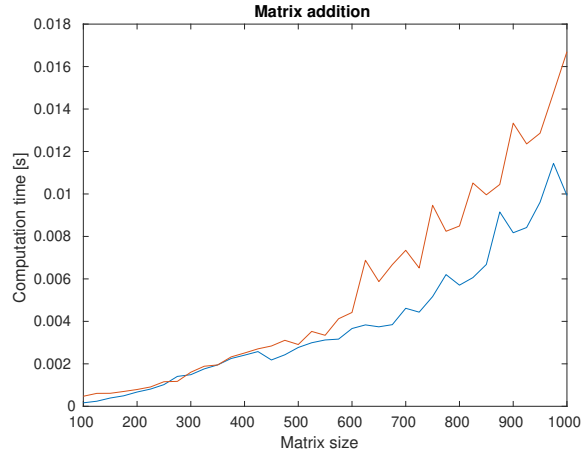


Figure 2: Computation time for matrix addition vs. matrix size with CPU (blue) and GPU(red).

For the matrix addition operation GPU implementation overall performance is comparable (or even worse) as CPU one. This is because additional resources are being spent for initializing memory tranfering data between host and device.

More complex algorithms, such as matrix multiplication and Gauss-Jordan elimination show different trends.




Figure 3: Computation time for matrix multiplication algorithm vs. matrix size with CPU (blue) and GPU(red). On the left panel: matrix size in range 100-1000, on the right: 10-100




Figure 4: Computation time for Gauss-Jordan algorithm vs. matrix size with CPU (blue) and GPU(red). On the left panel: matrix size in range 100-1000, on the right: 10-100

## 5.2 Nvprofiler analysis

An additional analysis was performed with the use of nvprof. Below are listed GPU activities for two example operations: matrix addition and Gauss-Jordan algorithm.

```
==31580== Profiling application: ./Toolkit gpu add normal
==31580== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	50.99%	2.4640us	3	821ns	736ns	960ns	[CUDA memcpy HtoD]
	32.45%	1.5680us	1	1.5680us	1.5680us	1.5680us	add kernel(double*, double*, double*, int, int)
	16.56%	800ns	1	800ns	800ns	800ns	[CUDA memcpy DtoH]
API calls:	99.58%	210.49ms	3	70.164ms	8.7840us	210.47ms	cudaMalloc
	0.14%	296.05us	1	296.05us	296.05us	296.05us	cuDeviceTotalMem
	0.11%	230.82us	96	2.4040us	281ns	91.470us	cuDeviceGetAttribute
	0.07%	143.55us	3	47.848us	10.162us	115.57us	cudaFree
	0.04%	84.044us	4	21.011us	10.948us	35.892us	cudaMemcpy
	0.04%	74.819us	1	74.819us	74.819us	74.819us	cudaLaunchKernel
	0.02%	52.023us	1	52.023us	52.023us	52.023us	cuDeviceGetName
	0.00%	4.9840us	1	4.9840us	4.9840us	4.9840us	cuDeviceGetPCIBusId
	0.00%	3.0990us	3	1.0330us	317ns	2.2710us	cuDeviceGetCount
	0.00%	1.3580us	2	679ns	327ns	1.0310us	cuDeviceGet
	0.00%	575ns	1	575ns	575ns	575ns	cuDeviceGetUuid

Figure 5: Profiling result for matrix addition operation: For matrix size n=10

```
==29809== Profiling application: ./Toolkit gpu inverse normal
==29809== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	39.22%	23.105us	20	1.1550us	1.1200us	1.4720us	reduce_row(double*, double*, int, int, bool)
	36.77%	21.664us	10	2.1660us	1.6320us	3.5200us	normalize_row(double*, double*, int, int)
	18.41%	10.848us	10	1.0840us	992ns	1.7280us	find_pivots(double*, double*, int, int)
	2.88%	1.6960us	2	848ns	736ns	960ns	[CUDA memcpy HtoD]
	2.72%	1.6000us	2	800ns	800ns	800ns	[CUDA memcpy DtoH]
API calls:	99.57%	226.25ms	2	113.13ms	7.6690us	226.24ms	cudaMalloc
	0.14%	322.26us	40	8.0560us	6.1850us	60.026us	cudaLaunchKernel
	0.12%	271.53us	1	271.53us	271.53us	271.53us	cuDeviceTotalMem
	0.07%	164.92us	96	1.7170us	158ns	70.664us	cuDeviceGetAttribute
	0.05%	105.95us	2	52.977us	13.227us	92.727us	cudaFree
	0.03%	67.484us	4	16.871us	8.3650us	27.253us	cudaMemcpy
	0.02%	39.975us	1	39.975us	39.975us	39.975us	cuDeviceGetName
	0.00%	5.8750us	1	5.8750us	5.8750us	5.8750us	cuDeviceGetPCIBusId
	0.00%	2.5980us	3	866ns	337ns	1.7540us	cuDeviceGetCount
	0.00%	1.1030us	2	551ns	219ns	884ns	cuDeviceGet
	0.00%	292ns	1	292ns	292ns	292ns	cuDeviceGetUuid

Figure 6: Profiling result for Gauss-Jordan algorithm: For matrix size n=10

It is visible that for the addition operation, the actual computation (by kernel) accounts for less than 50% of the GPU time, and the rest is spent to transfer data between host and device. For Gauss-Jordan elimination, on the other hand, most of the GPU time is spent by three kernels performing calculation, which means that we take more advantage of the parallel architecture.

## 6 Summary

A toolkit for linear algebra operations was implemented both on GPU using CUDA architecture, CPU code was also implemented as a reference point. The results (Fig. 2-4) show that GPU implementation provides significant speedup for algorithms such as matrix multiplication and Gauss-Jordan elimination, whereas is not very efficient for operations such as matrix addition/subtraction.